

# Databases

or: How I learned to stop worrying and love the deadlock

---

[Mario.Lassnig@cern.ch](mailto:Mario.Lassnig@cern.ch)

European Organization for Nuclear Research (CERN)

# About me

---

## ↪ Software Engineer

- ↪ Data Management & Analytics for the ATLAS Experiment, CERN, 2006-ongoing
- ↪ Automotive Navigation, AIT Vienna, 2004-2006
- ↪ Avionics for autonomous robots, Austrian Space Forum, 2008-ongoing

## ↪ Education

- ↪ Multivariate statistics and machine learning (PhD)
- ↪ Graph theory (Master's)
- ↪ Cryptography (Undergrad)

## ↪ Largest 24/7 database built yet

- ↪ 3+ billion rows
- ↪ 35'000 IOPS

# About this course

---

- ↪ For every topic
  - ↪ We will do some theory
  - ↪ We will do some hands-on exercises
- ↪ The contents from the session codes are complete and will give you the solution.
  - ↪ Don't blindly copy/paste, try to understand what's going on.
  - ↪ There'll be a challenge/exercise at the end, where you have to apply what you have learned.

Please interrupt me whenever necessary!

# Part I — Introduction

# CAP Theorem

It is impossible for a distributed computer system to simultaneously provide all three of the following guarantees [Brewer, 2000]

All clients always see the same data.  
**Consistency**

File systems, single-instance databases, ...

All clients can always read and write.  
**Availability**

Choose Two

Distributed databases

Web caching, DNS, ...

**Partition tolerance**  
The data can be split across the system.

# ACID and BASE

---

## ↪ ACID

- ↪ Atomicity All or nothing operations
- ↪ Consistency Always a valid state
- ↪ Isolation Operations can be serialised
- ↪ Durability Data is never lost

## ↪ BASE

- ↪ Basically available More often than not
- ↪ Soft state Data might be lost, but you will only know when you try to access it
- ↪ Eventually consistent Your query might return old data

# So what is this NoSQL thing anyway?

---

- ↪ Carlo Strozzi, 1998
- ↪ Term invented for a relational database without a SQL interface
- ↪ Term re-coined by *last.fm* in 2009
  - ↪ At an open-source distributed databases workshop
  - ↪ What they actually meant: how to deal with the exponential increase in storage requirements
- ↪ Promises of NoSQL databases?
  - ↪ Relational model might not map well to application's data structures
  - ↪ Improve productivity by using non-relational stores instead as application backend
- ↪ Improve performance for "web-scale" applications?
  - ↪ Remember the CAP theorem
  - ↪ There is no free lunch

# Types of databases

---

- ↪ Row stores
  - ↪ Oracle, PostgreSQL, MySQL, SQLite, ...
- ↪ Column stores
  - ↪ Hbase, Cassandra, Hypertable, Monet, ...
- ↪ Data structure stores
  - ↪ ElasticSearch, MongoDB, Redis, PostgreSQL, ...
- ↪ Many have overlapping concepts
- ↪ Get confused here:  
<http://nosql-database.org/>
- ↪ Key/Value stores
  - ↪ Dynamo, Riak, LevelDB, Kyoto, ...
- ↪ Graph stores
  - ↪ Neo4j, Titan, Hypergraph, ...
- ↪ Multimodel stores
  - ↪ ArangoDB, CortexDB, ...
- ↪ Object stores
  - ↪ Versant, Ceph, ...

Relational  
Non-relational

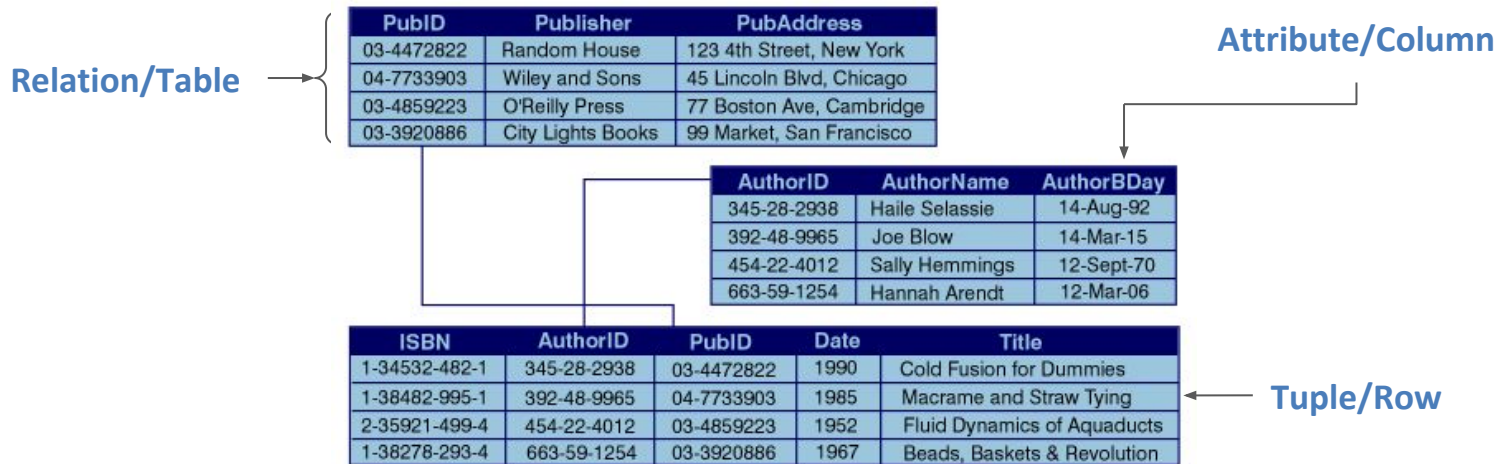


# Relational model

⇒ Proposed by Edgar F. Codd, 1969

⇒ **Concept**                      Relations                      Tuples                      Attributes

⇒ **Database**                      Table                      Row                      Column



<http://www.ibm.com/developerworks/library/x-matters8/relat.gif>

# Structured Query Language (SQL)

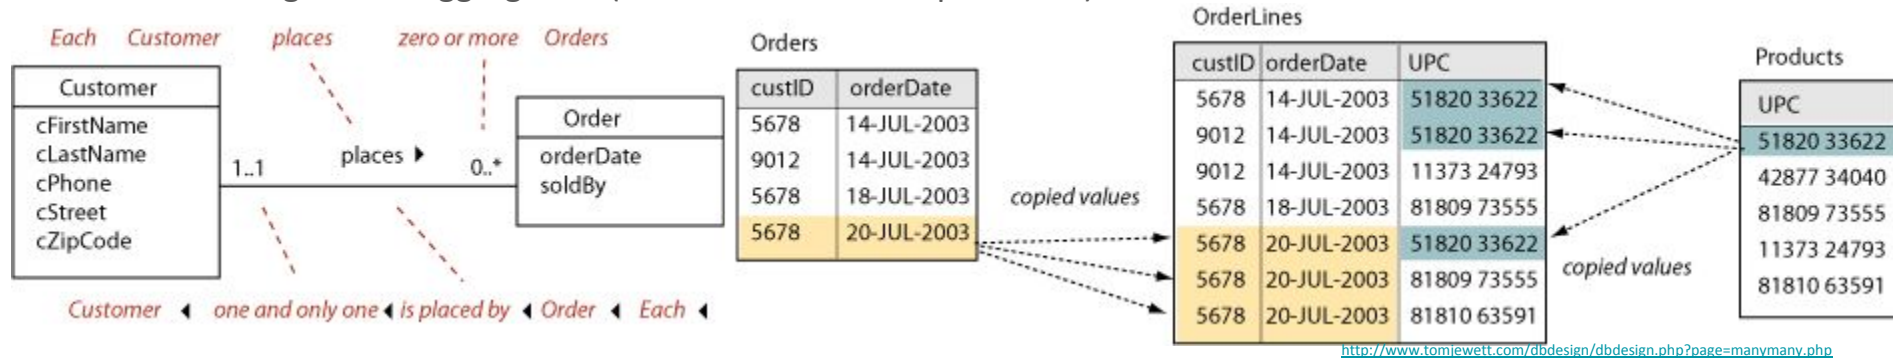
---

- ↪ Proposed by Edgar F. Codd, 1969
- ↪ Interaction with DBMS (Database Management System) using declarative language
- ↪ Developed from SEQUEL language an ANSI/ISO Standard in 1986
- ↪ Ess Que Ell? Sequel?
  - ↪ Oracle defines pronunciation as "sequel", then we have "My-S-Q-L", and PostgreSQL is "Postgres"
  - ↪ Don Chamberlin, one of the original authors of the SEQUEL language suggests to go with ISO: S-Q-L

```
CREATE TABLE table_name;  
SELECT column_name FROM table_name WHERE column_name = value;  
INSERT INTO table_name(column_name) VALUES (value);  
UPDATE table_name SET column_name = value;  
DELETE FROM table_name WHERE column_name = value;  
DROP TABLE table_name;
```

# Row Stores

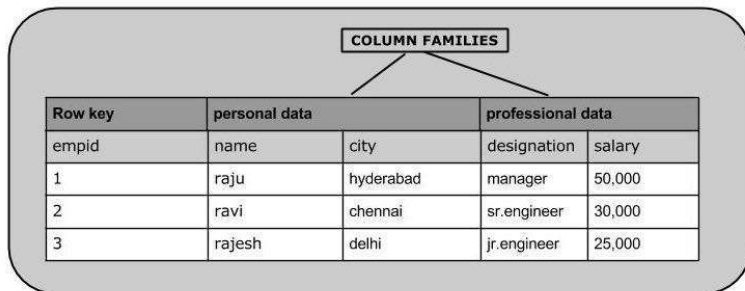
- ⇒ Your classic RDBMS
- ⇒ Physically stores data row-by-row
- ⇒ Easy joining of data between tables
  - ⇒ one-to-one
  - ⇒ one-to-many
  - ⇒ many-to-many
- ⇒ Normalization algorithms to reduce duplicate data and complexity
- ⇒ Not so good for aggregation (RDBMS vendors compete here)



<http://www.tomjewett.com/dbdesign/dbdesign.php?page=manymany.php>

# Column Stores


- ➡ Many applications do not need relations
- ➡ Row-based systems like traditional RDBMS are ill-suited for aggregation queries
  - ➡ Things like SUM/AVG of a column?
  - ➡ Needs to read blocks of columns, sometimes full rows, unnecessarily
- ➡ Physical layout of data column-wise instead
  - ➡ saves IO and improves compression, facilitates parallel IO
  - ➡ makes joins between columns harder
  - ➡ need to save on joins somehow
  - ➡ native support for column-families



# Key/Value Stores

---

- ↪ Hashmap for efficient insert and retrieval of data
  - ↪ You might know this as an associative array, or dictionary, or hashtable, ...
- ↪ Keys and values are usually bytestreams, but practically just strings
- ↪ Usually there are some performance guarantees, via options like
  - ↪ sorted keys
  - ↪ length restrictions
  - ↪ different hash functions
- ↪ Simple and easy to use
  - ↪ Either as compile-time libraries
  - ↪ Or as a server, usually via wrapped native protocols, or via REST
- ↪ First Key/Value store: dbm, 1979



# Data Structure Stores

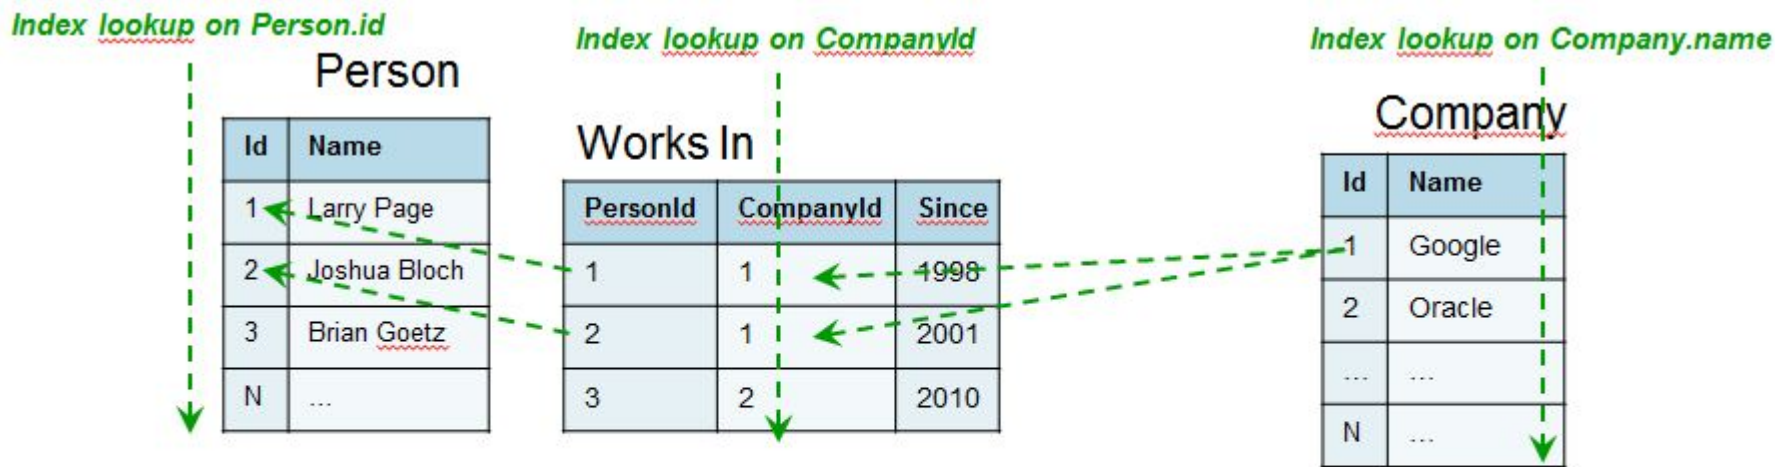
- ➡ Basically key/value stores, with the added twist that the stores knows something about the internal structure of the value
- ➡ Very easy to use as a backend for an application
- ➡ When people think NoSQL, this is usually what they mean



<http://docs.mongodb.org/v3.0/images/data-model-denormalized.png>

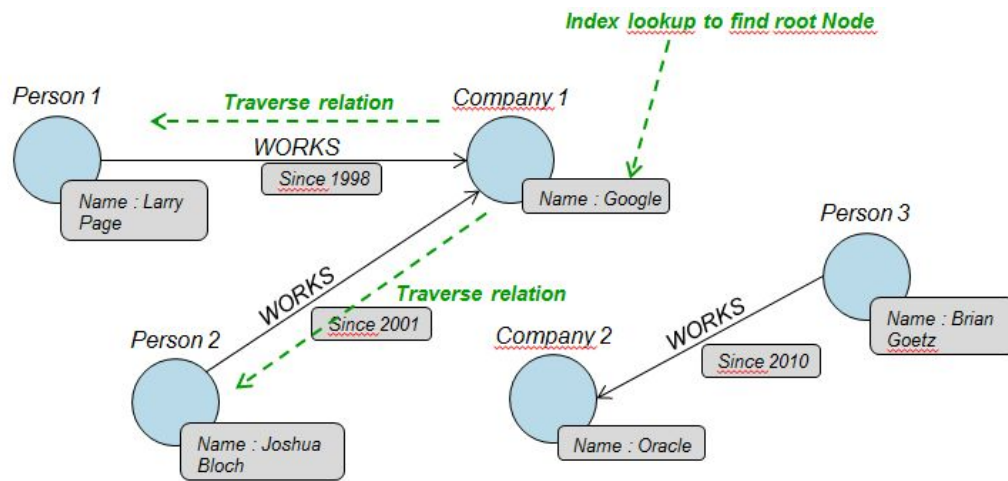
# Graph Stores

- ⇒ In the relational world, actual *n-to-n* relations are cumbersome to model (despite the name...)



# Graph Stores

- ➡ Make relations first-class citizens
- ➡ Physical layout optimised for distance between data points
- ➡ Leads to easy & fast traversal for the graph database engine





# Hands-on Session 1

---

- ↪ Create, read, update, delete data
- ↪ Using C/C++ and Python
- ↪ On
  - ↪ **PostgreSQL**      relational      row-based
  - ↪ **LevelDB**      non-relational      key/value
  - ↪ **Redis**      non-relational      data structure
  - ↪ **ElasticSearch**      non-relational      data structure
  - ↪ **Neo4j**      non-relational      graph

## Part II — Fun and profit

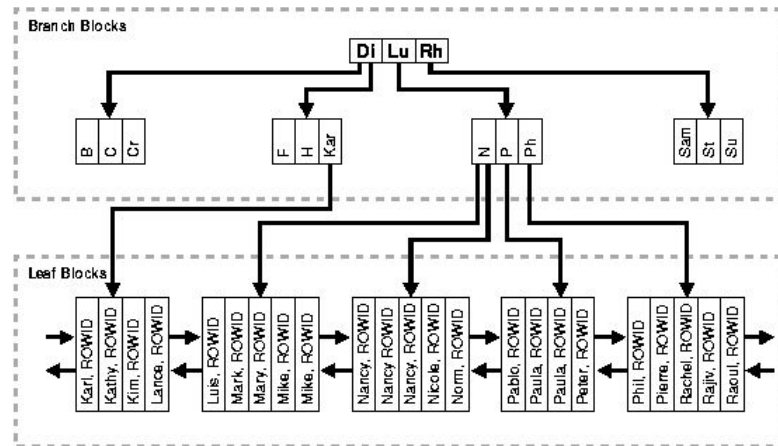
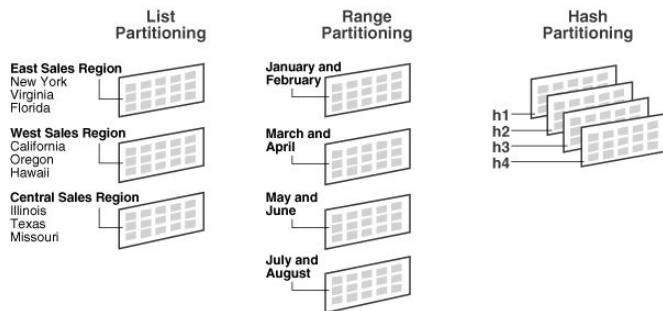
# Query plans

---

- ↪ IMHO the single most important thing when dealing with databases
- ↪ You want to avoid going to disk, to reduce the IOPS and CPU usage
- ↪ In order of excessiveness
  - ↪ FULL TABLE SCAN
  - ↪ PARTITION SCAN
  - ↪ INDEX RANGE SCAN
  - ↪ PARTITION INDEX RANGE SCAN
  - ↪ INDEX UNIQUE SCAN
  - ↪ PARTITION INDEX UNIQUE SCAN
- ↪ Not all FULL TABLE SCANS are bad
  - ↪ If you need to retrieve a lot of data, and it is indexed, you will get random IO on the disk. Much better to prefer serial scan (FULL, PARTITION) in such cases
  - ↪ If you data is low cardinality (few values, lots of rows) then indexes won't help and slow down INSERTs

# Query plans — How to optimise?

- Understand the EXPLAIN PLAN statement, then decide
- Partitions
  - Physical separation of data
  - Costly to introduce afterwards (usually requires downtime with schema migration)
- Indexes



# Transactional safety

↪ In multi-threaded environments concurrent access to the same data is likely

↪ Dirty Read

↪ Read data by uncommitted transaction

↪ Non-Repeatable Read

↪ Reads previously read data again, but it has changed in the meantime by another transaction

↪ Phantom Read

↪ Repeated query of the same conditions yields different results due to intermediate other transaction

↪ Different transaction isolation levels provide safeguards

↪ By locking of rows, or sometimes even tables, and thus making other transactions wait

↪ The more you lock, the slower you are

↪ Can lead to deadlocks if the developer is careless — always lock rows in the same order!

Isolation Level	Dirty Read	Nonrepeatable Read	Phantom Read
Read uncommitted	Possible	Possible	Possible
Read committed	Not possible	Possible	Possible
Repeatable read	Not possible	Not possible	Possible
Serializable	Not possible	Not possible	Not possible

# Sparse metadata

---

- ↪ Think "tagging"/"labelling" datasets — Difficult problem in relational model
  - ↪ Keep extra columns or implement a relational key/value store?
  - ↪ Extra columns are bad for physical disk layout
  - ↪ Relational key/value store requires lots of joins — not good for cache and CPU
  
- ↪ Will you ever search on metadata?
  - ↪ **No**            Store the metadata as a JSON-encoded string in a single column
  - ↪ **Yes**            Many different kinds of metadata?
    - **No**            Maintain a separate metadata table, with pre-created columns
    - **Yes**           Use the built-in JSON support of Oracle or PostgreSQL  
or alternatively use a non-relational database and keep the data in sync
  
- ↪ There is some really bad advice on StackOverflow promoting a generic key/value model for relational databases — Caveat emptor!

# Competitive locking

---

- ↪ Many applications have the following use case
  - ↪ Many processes write something in a queue — the "backlog" of things to do
  - ↪ Many processes read from that queue — and work it off in parallel
- ↪ Scheduling problem
  - ↪ Do things in order? Prioritise certain things?
  - ↪ How to avoid that multiple workers process the same things?
  - ↪ Repeat after me: a database is not a queuing system
- ↪ However, sometimes a DB is all you have... And there are two potential solutions
- ↪ Database-level (row read lock, easy):
  - ↪ `BEGIN;        SELECT row FOR UPDATE;                COMMIT/ROLLBACK;`
- ↪ Application-level (no lock, complex)
  - ↪ When selecting work, compute the row-hash, convert to integer, take modulo #workers
  - ↪ Then only process rows that match the worker-id

## Part III — Survival



# Distributed transactions

- ↪ Sometimes you really need two different data stores
- ↪ Sometimes you need to be consistent between both
- ↪ Consensus protocols are needed
  - ↪ Two-phase commit
  - ↪ Paxos
- ↪ Needs operational support by DBMS
  - ↪ pending writes / write-ahead log (WAL)
- ↪ But you still have to code it in the application
  - ↪ This is where you will lose performance

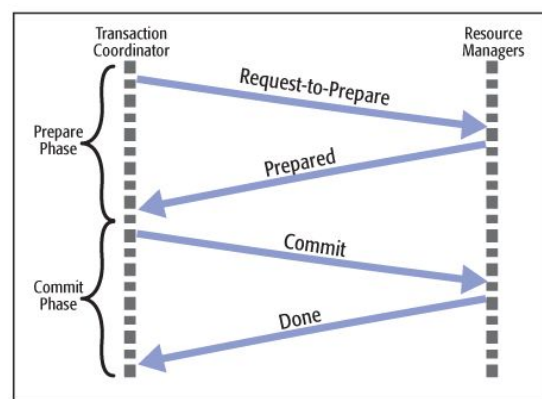



Figure 1 • The two-phase commit protocol

# SQL Injection



# SQL Injection

---



<https://hackaday.com/files.wordpress.com/2014/04/18mpenleoksg8jpg.jpg?w=636>

# Application level woes

---

- ↪ Handling sessions — this will be your major source of pain
  - ↪ Connections, Sessions, Transactions (and how they are differently named/treated by each DBMS)
- ↪ Most databases only have a limited number of available connections, or sessions
- ↪ Some pay with CPU for logons, others with RAM for sessions, etc...
- ↪ E.g., if you have to channel 100 concurrent transactions across 10 connections
  - ↪ SessionPooling/QueuePooling — every language/database has its own idea
  - ↪ Takeaway message: Use an abstraction, and don't code it y ourself
- ↪ Python            SQLAlchemy, Django
  - ↪ Also comes with Object-Relational Mapper
  - ↪ Automatically translates Python objects into the relational model
- ↪ C++                CodeSynthesisODB
  - ↪ Also supports BOOST datatypes!

## Part IV — The Challenge

# Challenge description

---

- ↪ Write a Twitter clone before time runs out
- ↪ Choose any database you like, the VMs are preconfigured for all of them
- ↪ Stick to the following UX — you will need to write four programs
  - ↪ `Inserter`                      periodically inserts new random tweet into the database
  - ↪ `Latest`                        periodically print the latest 10 tweets
  - ↪ `Random`                        periodically print random 5 tweets
  - ↪ `Stats`                         periodically print statistics about the database content
    - How many tweets were added in the last minute by each user and overall (insertion rate)
    - How often a given tweet was displayed (popularity of a particular twwet)
  - ↪ Start 10 inserter, 10 latest, 10 random, 1 stat
- ↪ Use this random tweet/sentence generator (`pip install loremipsum`)
  - ↪ `import loremipsum`                      `loremipsum.generate_sentence()`
- ↪ When stuck, ask me — when done, show me! Good luck and have fun!

# Databases

or: How I learned to stop worrying and love the deadlock

---

[Mario.Lassnig@cern.ch](mailto:Mario.Lassnig@cern.ch)

European Organization for Nuclear Research (CERN)

